

# Different ways to improve python code efficiency

Pierre-Yves Dupont

## Simple example

# Fibonacci number

## fibonacci.py

```
1 def fibo(n):  
2     if n == 0 or n == 1:  
3         return n  
4     return fibo(n - 2) + fibo(n - 1)
```

Pure python version: fibo(38) computed in 15s

# Cython

- `apt-get install cython`

- apt-get install cython
- programming language based on python (Pyrex)

- apt-get install cython
- programming language based on python (Pyrex)
- simplify development of C extensions for Python

- apt-get install cython
- programming language based on python (Pyrex)
- simplify development of C extensions for Python
- possible to transform python code in C code



- apt-get install cython
- programming language based on python (Pyrex)
- simplify development of C extensions for Python
- possible to transform python code in C code
- used in Scipy or SAGE

## Cython on fibo.py

```
1 #build fibo.c file
2 cython fibo.py
3 #build fibo.so library
4 gcc fibo.c -o fibo.so -shared -pthread -fPIC \
5     -fwrapv -O2 -Wall -fno-strict-aliasing $(pkg-config python --cflags)
```

## Using fibo.so

```
1 import fibo
2 print fibo.fibo(38)
```

Simple Cython version: fibo(38) computed in 10s (33%)

# More improvement

## fibonacci.pyx

```
1 cdef int fibo(int n):  
2     if n == 0 or n == 1:  
3         return n  
4     return fibo(n-2) + fibo(n-1)
```

## Cython compiling

```
1 cython fibo.py  
2 gcc fibo.c -o fibo.so -shared -pthread -fPIC \  
3     -fwrapv -O2 -Wall -fno-strict-aliasing $(pkg-config python --cflags)
```

Cython/pyx version: fibo(38) computed in 5s (66.7%)

# Standalone executable

```
1 cython --embed test_fibo.py -o fibo.c
2 gcc fibo.c -o fibo $(pkg-config python --cflags) -lpthon2.7
3 ./fibo
```

See more on <http://docs.cython.org/index.html>

# Pure C version

## fibonacci.c

```
1 #include <stdio.h>
2 #include <stdlib.h>
3
4 long fibo(long n){
5     if(n == 0 || n == 1){return n;}
6     return fibo(n - 2) + fibo(n - 1);
7 }
8
9 void main(int argc, char* argv[]){
10     printf("%ld\n", fibo(38));
11 }
```

C version: fibo(38) computed in 0.5s

## How to extend python with pure C libraries

- Simplified Wrapper and Interface Generator

# SWIG

- Simplified Wrapper and Interface Generator
- Tool used to connect C/C++ libraries to programs written in Python, Perl, Ruby, R, Java. . .



# SWIG

- Simplified Wrapper and Interface Generator
- Tool used to connect C/C++ libraries to programs written in Python, Perl, Ruby, R, Java. . .
- `apt-get install swig`

# Connect fibo\_pure\_c.c to Python

## Interface file fibo\_pure\_c.i

```
1 %module fibo_pure_c
2 %{
3 extern long fibo(long n);
4 %}
5 extern long fibo(long n);
```

## Wrapper generation

```
1 #generate fibo_pure_c_wrap.c
2 swig -python fibo_pure_c.i
3 #build the library
4 gcc -c fibo_pure_c.c fibo_pure_c_wrap.c \
5     -I/usr/include/python2.7
6 ld -shared fibo_pure_c.o fibo_pure_c_wrap.o -o _fibo_pure_c.so
```

## Use the C library in python

```
1 import fibonacci as fib  
2 print fib.fib(38)
```

Python with pure C library: computation of fibo(38) in 0.5s

Another problem: evaluation of  $\pi$

Pi

$$\pi = \int_0^1 f(x) \, dx, \text{ with } f(x) = \frac{4}{1+x^2} \quad (1)$$

$$\pi = \frac{1}{n} \sum_{i=1}^n f(x_i), \text{ with } x_i = \frac{i - \frac{1}{2}}{n} \text{ for } i = 1, \dots, n \quad (2)$$

```

1 import sys
2 PI = 3.141592653589793
3
4 def f(a):
5     return 4.0/(1.0+a**2)
6
7 def main():
8     while(1):
9         n = raw_input("Enter the number of intervals: (0 quits)\n")
10        try:
11            n=int(n)
12        except ValueError:
13            return 2
14        if n == 0: break
15        #LOOP TO PARALLELIZE
16        h = 1.0/n
17        sum = 0.0
18        for i in range(1,n):
19            x = h*(i - 0.5)
20            sum += f(x)
21        pi = h * sum
22        #END
23        sys.stdout.write("pi is approximatly: %.16f Error is: %.16f \n"\
24        % (pi, abs(pi-PI)));

```

$n = 10^8$  pi evaluated in 25s,  $n = 10^9$  crash

# Weave

```
1 import sys
2 from scipy import weave
3 from scipy.weave import converters
4 import time
5 PI = 3.141592653589793
6
7 code="""
8 int i;
9 double x;
10 double sum = 0.0;
11 for(i = 1;i <= n; i++) {
12     x = h*((double)i-(double)0.5);
13     sum += (double)4.0/((double)1.0+(x*x));
14 }
15 return_val = sum;
16 """
17 vars = "h n".split()
18 ...
```

# Weave

```
1 ...
2 def main():
3     while(1):
4         n = raw_input("Enter the number of intervals: (0 quits)\n")
5         try:
6             n=int(n)
7         except ValueError:
8             return 2
9         if n == 0: break
10        #LOOP TO PARALLELIZE
11        h = 1.0/n
12        sum = float(0.0)
13        tps = time.time()
14        sum = weave.inline(code, vars,
15                           type_converters = converters.blitz,
16                           compiler = 'gcc')
17        print time.time() - tps
18        pi = h * sum
19        #END
20        sys.stdout.write("pi is approximatly: %.16f Error is: %.16f \n"\
21                          % (pi, abs(pi-PI)));
```

$n = 10^8$  pi evaluated in 0.3s,  $n = 10^9$  pi evaluated in 4s



## MPI and OpenMP

# Parallel computing with MPI and OpenMP

## MPI

- Message Passing Interface

# Parallel computing with MPI and OpenMP

## MPI

- Message Passing Interface
- Fortran 77, 90, 95 and C/C++

# Parallel computing with MPI and OpenMP

## MPI

- Message Passing Interface
- Fortran 77, 90, 95 and C/C++
- Interfaces for C#, Java...

# Parallel computing with MPI and OpenMP

## MPI

- Message Passing Interface
- Fortran 77, 90, 95 and C/C++
- Interfaces for C#, Java...
- Parallel machines and on workstation clusters

# Parallel computing with MPI and OpenMP

## OpenMP

- Open MultiProcessing

# Parallel computing with MPI and OpenMP

## OpenMP

- Open MultiProcessing
- Fortran 90, 95 and C/C++

# Parallel computing with MPI and OpenMP

## OpenMP

- Open MultiProcessing
- Fortran 90, 95 and C/C++
- Shared memory multiprocessing programming



# Parallel computing with MPI and OpenMP

## OpenMP

- Open MultiProcessing
- Fortran 90, 95 and C/C++
- Shared memory multiprocessing programming
- Most of the processor architectures compatible with MPI

# Parallel computing with MPI and OpenMP

## OpenMP

- Open MultiProcessing
- Fortran 90, 95 and C/C++
- Shared memory multiprocessing programming
- Most of the processor architectures compatible with MPI
- Based on *pragmas* (compiler-specific preprocessor directives)

# Simple C version

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <math.h>
4 #define PI 3.1415926535897932384626433832795029L
5
6 double f(double a) {
7     return (double)4.0/((double)1.0+(a*a));
8 }
9 ...
```

# Simple C version

```
1 int main(int argc, char* argv[])
2 {
3     int n, i;
4     double h, pi, sum, x ;
5     for(;;) {
6         printf("Enter the number of intervals: (0 quits)");
7         if(!scanf("%lu",&n)){return 2;}
8         if(n ==0)
9             break;
10        //LOOP TO PARALLELIZE
11        h = ((double)1.0)/((double)n;
12        sum = 0.0;
13        for(i =1;i<=n;i++) {
14            x = h*((double)i-(double)0.5);
15            sum += f(x);
16        }
17        pi = h*sum;
18        //END
19        printf("pi is approximatly: %.16f Error is: %.16f \n",\
20            pi, fabs(pi-PI));
21    }
22    return EXIT_SUCCESS ;
23 }
24 $n = 10^9$ pi evaluated in 4sec
```

# Pi evaluation using MPI

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <math.h>
4 #include <mpi.h>
5 #define PI 3.1415926535897932384626433832795029L
6
7 double f(double a) {
8     return (double)4.0/((double)1.0+(a*a));
9 }
10 ...
```

# Pi evaluation using MPI

```
1 int main(int argc, char* argv[])
2 {
3     int n, i;
4     double h, pi, sum, x ;
5
6     double mypi;
7     int myid, numprocs, islave;
8     MPI_Status status;
9     MPI_Init(&argc, &argv);
10    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
11    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
12    n=0;
13    ...
```

```

1  for(;;) {
2      if (myid == 0){
3          printf("Enter the number of intervals: (0 quits)\n");
4          if(!scanf("%d",&n)){return 2;}
5      }
6      MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
7      if(n ==0)
8          break;
9      //LOOP TO PARALLELIZE
10     h = ((double)1.0)/(double)n;
11     sum = 0.0;
12     for(i = myid + 1; i <= n; i += numprocs){
13         x = h * ((double)i - (double)0.5);
14         sum += f(x);
15     }
16     mypi = h*sum;
17     MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
18     //END
19     if(myid == 0){
20         printf("pi is approximatly: %.16f Error is: %.16f \n",\
21             pi, fabs(pi-PI));
22     }
23 }
24 MPI_Finalize();
25 return EXIT_SUCCESS ;
26 }

```

## Compilation and Run

```
mpicc -g -O2 pi_mpi.c -o pi_mpi #mpi_c++ exists  
mpirun -np 12 -hostfile hostfile ./pi_mpi
```

- Different compiler, some differences with standard gnu compiler



## Compilation and Run

```
mpicc -g -O2 pi_mpi.c -o pi_mpi #mpi_c++ exists  
mpirun -np 12 -hostfile hostfile ./pi_mpi
```

- Different compiler, some differences with standard gnu compiler
- The code is completely parallelized

## Compilation and Run

```
mpicc -g -O2 pi_mpi.c -o pi_mpi #mpi_c++ exists  
mpirun -np 12 -hostfile hostfile ./pi_mpi
```

- Different compiler, some differences with standard gnu compiler
- The code is completely parallelized
- A lot of differences between simple C code and MPI code

## Compilation and Run

```
mpicc -g -O2 pi_mpi.c -o pi_mpi #mpi_c++ exists  
mpirun -np 12 -hostfile hostfile ./pi_mpi
```

- Different compiler, some differences with standard gnu compiler
- The code is completely parallelized
- A lot of differences between simple C code and MPI code
- $n = 10^9$  pi evaluated in 2s (12 threads)

# Pi evaluation using OpenMP

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <math.h>
4 #include <omp.h>
5 #define PI 3.1415926535897932384626433832795029L
6
7 double f(double a) {
8     return (double)4.0/((double)1.0+(a*a));
9 }
10 ...
```

```

1 int main(int argc, char* argv[]){
2     int n, i;
3     double h, pi, sum, x ;
4     for(;;) {
5         printf("Enter the number of intervals: (0 quits)\n");
6         if(!scanf("%u",&n)){return 2;}
7         if(n ==0)
8             break;
9         //LOOP TO PARALLELIZE
10        h = ((double)1.0)/(double)n;
11        sum = 0.0;
12        #pragma omp parallel for num_threads(12) private(i,x) reduction(+:sum)
13        for(i =1;i<=n;i++) {
14            x = h*((double)i-(double)0.5);
15            //#pragma omp critical
16            sum += f(x);
17        }
18        pi = h*sum;
19        //END
20        printf("pi is approximatly: %.16f Error is: %.16f \n",\
21            pi, fabs(pi-PI));
22    }
23    return EXIT_SUCCESS ;
24 }

```

## Compilation and Run

```
gcc -fopenmp -Wall -O2 pi_openmp.c -o pi_openmp #g++ works fine  
./pi_openmp
```

- Standard compiler, use openmp library

## Compilation and Run

```
gcc -fopenmp -Wall -O2 pi_openmp.c -o pi_openmp #g++ works fine  
./pi_openmp
```

- Standard compiler, use openmp library
- Only part of code (loops) are parallelized

## Compilation and Run

```
gcc -fopenmp -Wall -O2 pi_openmp.c -o pi_openmp #g++ works fine  
./pi_openmp
```

- Standard compiler, use openmp library
- Only part of code (loops) are parallelized
- Really close to original C code



## Compilation and Run

```
gcc -fopenmp -Wall -O2 pi_openmp.c -o pi_openmp #g++ works fine  
./pi_openmp
```

- Standard compiler, use openmp library
- Only part of code (loops) are parallelized
- Really close to original C code
- $n = 10^9$  pi evaluated in 3s (12 threads)

# Pi evaluation using OpenMP and MPI

In MPI version of code, insertion of OMP pragma

```
1  h = ((double)1.0)/((double)n;
2  sum = 0.0;
3  #pragma omp parallel for reduction(+:sum) private(i,x)
4  for(i = myid + 1; i <= n; i += numprocs){
5      x = h * ((double)i - (double)0.5);
6      sum += f(x);
7  }
8  mypi = h*sum;
9  MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
```

## Compilation and Run

```
mpicc -g -O2 -fopenmp pi_mpi.c -o pi_mpi  
mpirun -np 12 -hostfile hostfile ./pi_mpi
```

- MPI compiler

## Compilation and Run

```
mpicc -g -O2 -fopenmp pi_mpi.c -o pi_mpi  
mpirun -np 12 -hostfile hostfile ./pi_mpi
```

- MPI compiler
- MPI used to share the computation on different nodes

## Compilation and Run

```
mpicc -g -O2 -fopenmp pi_mpi.c -o pi_mpi  
mpirun -np 12 -hostfile hostfile ./pi_mpi
```

- MPI compiler
- MPI used to share the computation on different nodes
- OpenMP used to make threads within a node

## Compilation and Run

```
mpicc -g -O2 -fopenmp pi_mpi.c -o pi_mpi  
mpirun -np 12 -hostfile hostfile ./pi_mpi
```

- MPI compiler
- MPI used to share the computation on different nodes
- OpenMP used to make threads within a node
- $n = 10^9$  pi evaluated in 2s (6 nodes, 2 threads)

# Inlining parallelized C/C++ code in python

```
import sys
from scipy import weave
from scipy.weave import converters
import time

PI = 3.141592653589793

code="""
int i;
double x;
double sum = 0.0;
#pragma omp parallel for private(i,x) reduction(+:sum)
for(i = 1;i <= n; i++) {
    x = h*((double)i-(double)0.5);
    sum += (double)4.0/((double)1.0+(x*x));
}
return_val = sum;
"""

vars = "h n".split()
weave_omp = \
{
    'headers': ['<omp.h>'],
    'extra_compile_args': ['-fopenmp'],
    'extra_link_args': ['-lgomp']
}
```

# Inlining parallelized C/C++ code in python

```
def f(a):  
    return 4.0/(1.0+a**2)  
  
def main():  
    while(1):  
        n = raw_input("Enter the number of intervals: (0 quits)\n")  
        try:  
            n=int(n)  
        except ValueError:  
            return 2  
        if n == 0: break  
        #LOOP TO PARALLIZE  
        h = 1.0/n  
        sum = float(0.0)  
        tps = time.time()  
        sum = weave.inline(code, vars,  
                            type_converters = converters.blitz,  
                            compiler = 'gcc', **weave_omp)  
        print time.time() - tps  
        pi = h * sum  
        #END  
        sys.stdout.write("pi is approximatly: %.16f Error is: %.16f \n" %  
                        (pi, abs(pi-PI)));
```

$n = 10^9$  pi evaluated in 1.3s (12 threads)



# Comparison of the efficiency of the different methods

slow	430 seconds
numeric	2.76 seconds
Cython pyrex	2.55 seconds
fortran77	2.53 seconds
fortran90	0.60 seconds
fortran95	0.59 seconds
C	2.20 seconds